Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE 100%

Avoid Over-fitting Using Regularization

LATEST SUBMISSION GRADE

100%

1. What are the two types of regularization techniques that you learned in this course? 1 / 1 point

- [ ] Batch Normalization
- [x] Dropout Regularization

Correct
Correct, this was one of the techniques discussed.

- [x] Weight Regularization

Correct
Correct, this was one of the techniques discussed.

2. Let's say you have trained a Neural Network model with 10,000 training examples and 2,500 validation examples. At the end of training, you see a training accuracy of 90% and validation accuracy of 93%. Is it a case of the model over-fitting the training data? 1 / 1 point

- (•) No
- ( ) Yes

Correct
Correct. The validation accuracy is higher than the training accuracy. Therefore, it is not a case of model over-fitting the training data.

3. Let's say you have 4 classes in your data-set represented as 0, 1, 2 or 3. If you were to one-hot-encode a label with value 2, which one would be the encoded label? 1 / 1 point

- ( ) [1, 1, 0, 1]
- (•) [0, 0, 1, 0]
- ( ) [0, 1, 0, 0]

Correct
Correct.

4. Take a look at the following plot. 1 / 1 point

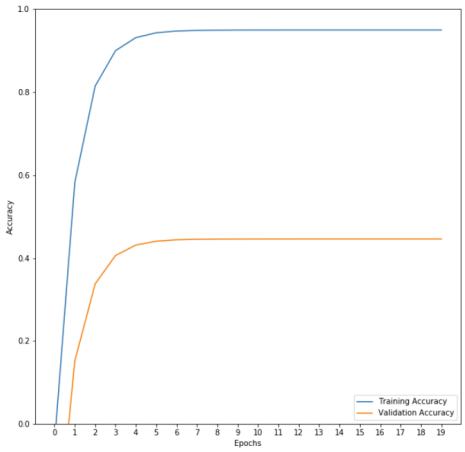

As shown, the blue plot represents the training accuracy and the orange plot represents the validation accuracy. Is this a case of over-fitting?

- ( ) No
- (•) Yes

Correct
Correct.

5. Which function returns the history object which has training metrics like training accuracy and loss? 1 / 1 point

- ( ) model.compile()
- (•) model.fit()

Correct
Correct.

6. What is the parameter that we set to apply L2 weight regularization in our model's dense layers? 1 / 1 point

- ( ) kernel_initializer
- ( ) l2
- (•) kernel_regularizer

Correct
Correct

Activate Windows